



Eucalyptus

Eucalyptus: An Open-source Infrastructure for Cloud Computing

Shashi Mysore
Eucalyptus Systems Inc.
www.eucalyptus.com

Open-source cloud computing infrastructure

- **Implements private and hybrid clouds**
- **Uses local machines and data without modification**
- **Transforms your resources into a cloud computing system, such as Amazon EC2**

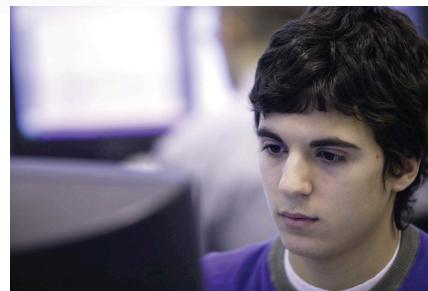


Company Facts



- **Creators of the leading open-source private cloud platform**
- **Incorporated January 2009**
- **Headquartered in Santa Barbara, CA**
- **Funded by NEA, Benchmark Capital, and BV Capital**
- **Born from the EUCLYPTUS cloud computing research project from U.C. Santa Barbara (released May 2008)**

What is a cloud?



SLAs

Web Services



Virtualization

Open-source Cloud Infrastructure



- **Idea: Develop an open-source, freely available cloud platform for commodity hardware and software environments**
 - Stimulate interest and build community knowledge
 - Quickly identify useful innovations
 - Act to dampen the “hype”
- **First-principles cloud implementation**
 - Not a refactoring of previously developed technology

What's in a name?



- **Elastic Utility Computing Architecture Linking Your Programs To Useful Systems**
- **Web services based implementation of elastic/utility/cloud computing infrastructure**
 - Linux image hosting ala Amazon
- **How do we know if it is a cloud?**
 - Try and emulate an existing cloud: Amazon AWS
- **Functions as a software overlay**
 - Existing installation should not be violated (too much)
- **Focus on installation and maintenance**
 - “*System Administrators are people too.*”



Goals for Eucalyptus

- **Foster greater understanding and uptake of cloud computing**
 - Provide a vehicle for extending what is known about the utility model of computing
- **Experimentation vehicle prior to buying commercial services**
 - Provide development, debugging, and “tech preview” platform for Public Clouds
- **Homogenize local IT environment with Public Clouds**
 - AWS functionality locally makes moving using Amazon AWS easier, cheaper, and more sustainable
- **Provide a basic software development platform for the open source community**
 - E.g. the “Linux Experience”
- **Not designed as a replacement technology for AWS or any other Public Cloud service**



Requirements

- **Implement cloud abstractions and semantics**
 - Must be a cloud (inarguably)
- **Simple**
 - Must be transparent and easy to understand
- **Scalable**
 - Interesting effects are observed at scale (e.g. not an SDK)
- **Extensible**
 - Must promote experimentation
- **Non-invasive**
 - Must not violate local control policies
- **System Portable**
 - Must not mandate a system software stack change
- **Configurable**
 - Must be able to run in the maximal number of settings
- **Easy**
 - To distribute, install, secure, and maintain

Open-source Cloud Anatomy



- **Extensibility**
 - Simple architecture and open internal APIs
- **Client-side interface**
 - Amazon's AWS interface and functionality (familiar and testable)
- **Networking**
 - Virtual private network per cloud
 - Must function as an overlay => cannot supplant local networking
- **Security**
 - Must be compatible with local security policies
- **Packaging, installation, maintenance**
 - system administration staff is an important constituency for uptake

The Elements of Cloud Style



- **SaaS (Software as a Service)**
 - Applications exporting network-facing user interfaces
 - User transfers data to the cloud
- **PaaS (Platform as a Service)**
 - Program or scripting runtime exports network-facing interfaces
 - Internal platform services available
 - User transfers program code and data to the cloud
- **IaaS (Infrastructure as a Service)**
 - Resource provisioning services export network-facing interfaces
 - Internal platform services available
 - User transfers code, data, and environment to the cloud



Why IaaS?

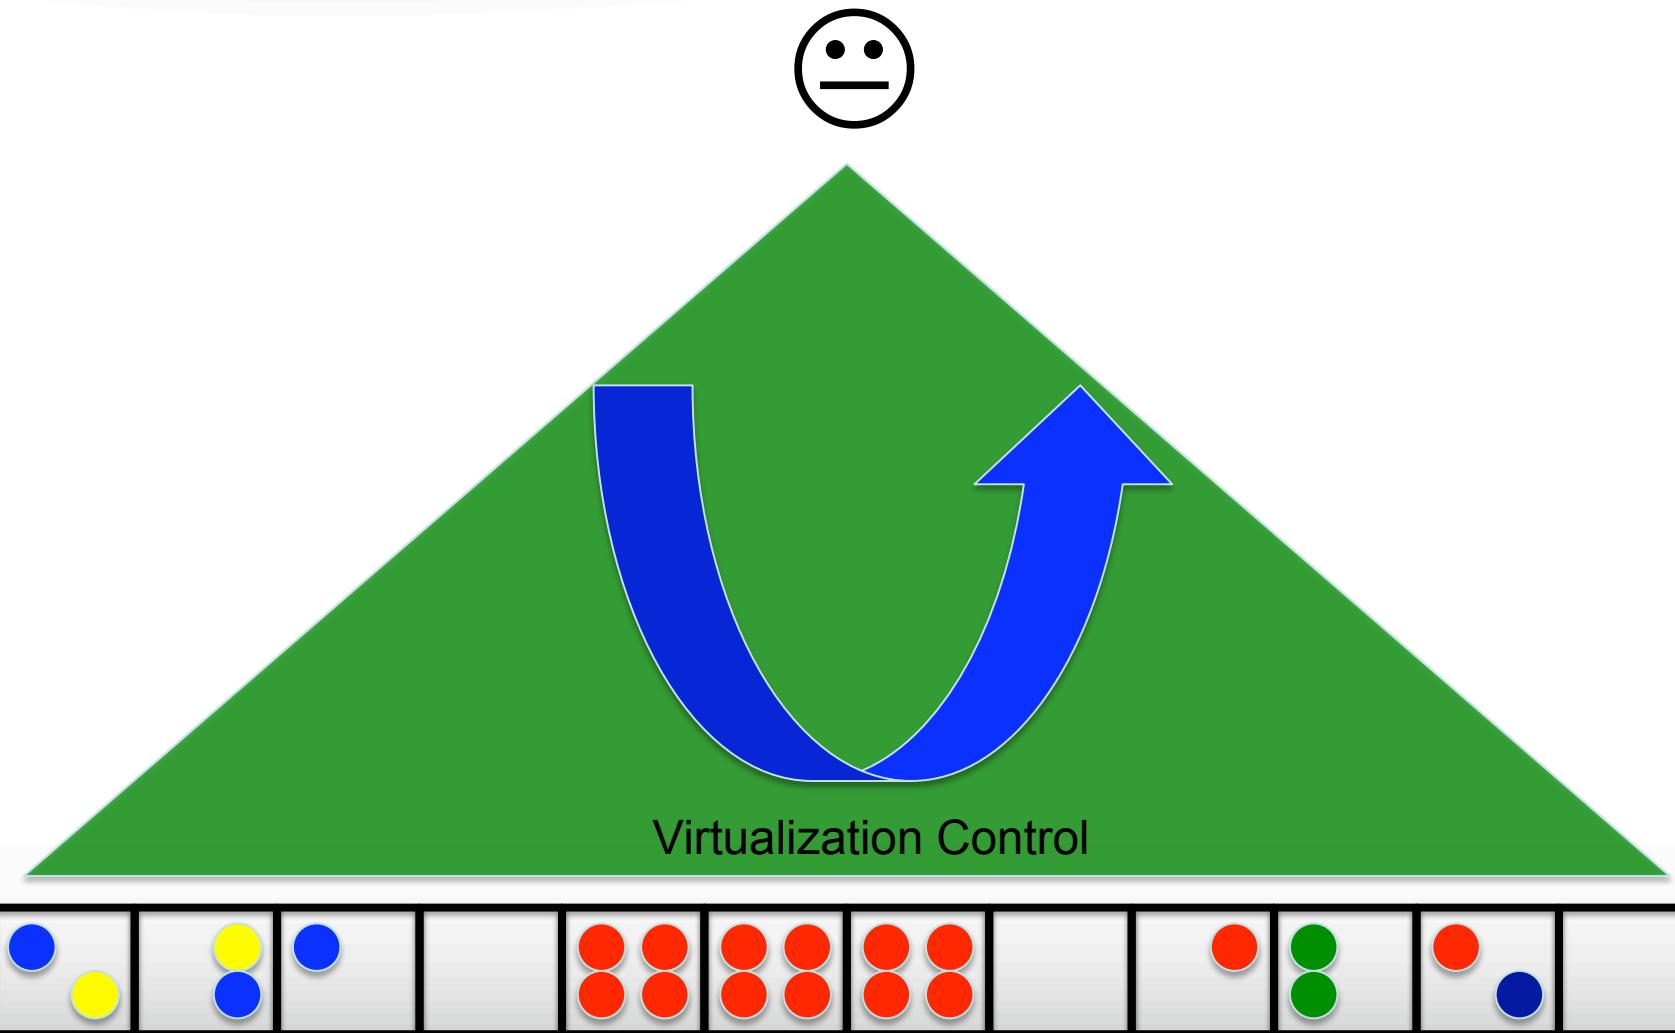
- **Applications are often multi-technology**
 - System “images” for different technologies can be combined
 - Multiple language technologies at different revision levels
- **Legacy support**
 - System images that mimic bare metal deployments can be used
 - Legacies are archived with the environment necessary to run them
- **Transparency**
 - Debugging and performance tuning can go down to the hypervisor
- **QoS containers**
 - QoS is implemented in the infrastructure today => familiar
- **Anti-lock in**
 - If clouds fail, a return path to bare metal is available

Clouds and Virtualization

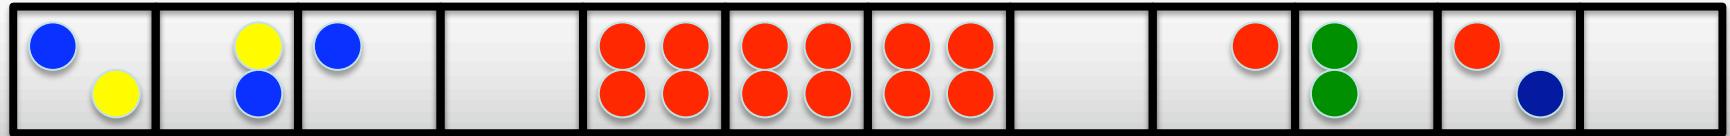
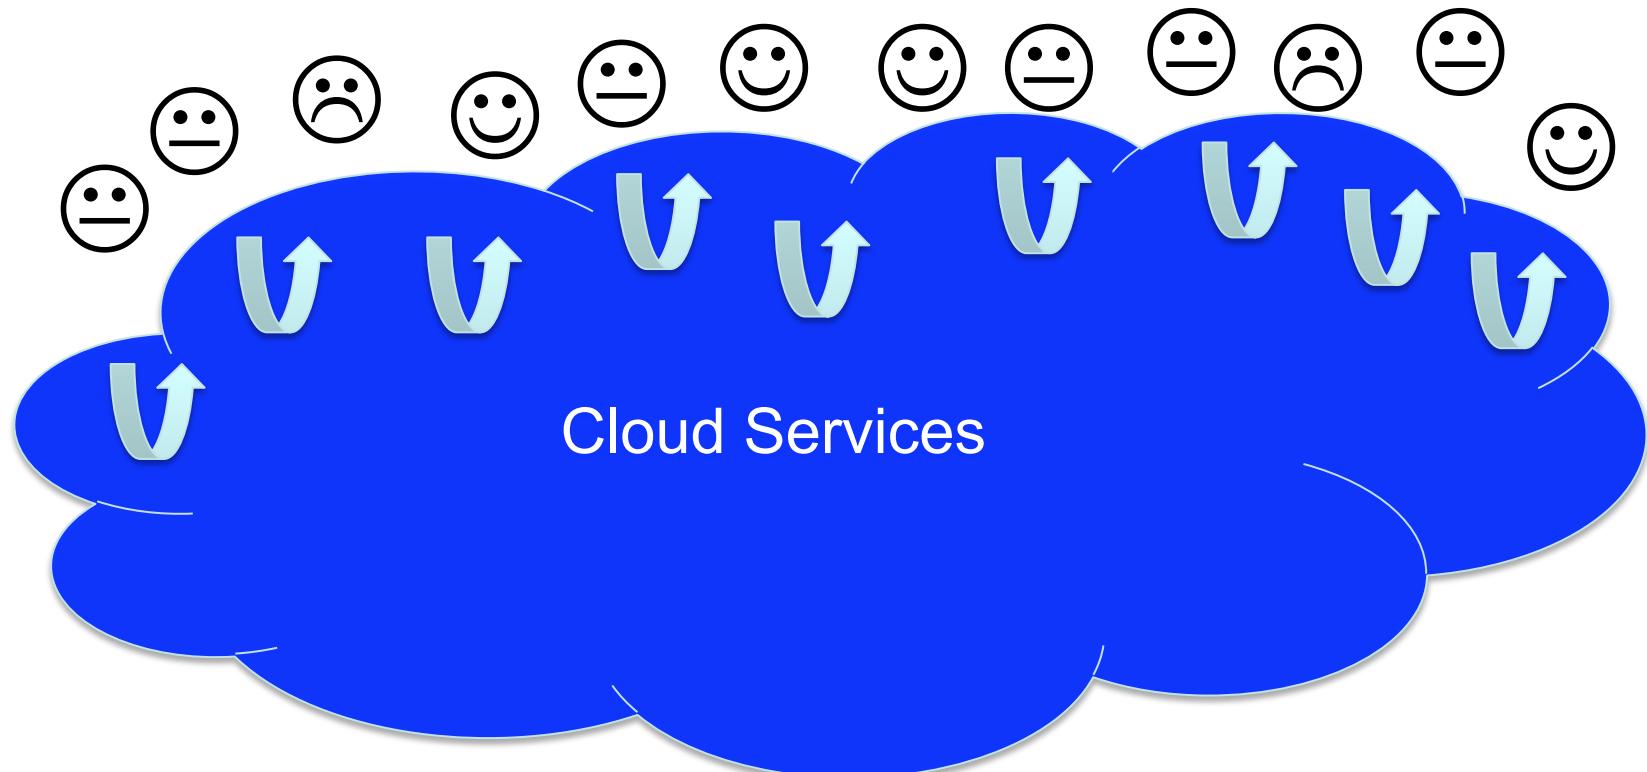


- ***Isn't a virtualized data center a private cloud?***
 - “I’m running Xen in my data center – I’m running a private cloud.”
- **Web service interfaces to hypervisors are useful**
 - Clouds require transactional operations across multiple hypervisors
- **Data center virtualization platforms are synchronous**
 - Data center administrators typically requires sequential consistency

Datacenter Virtualization



A Cloud



Cloud QoS Factors



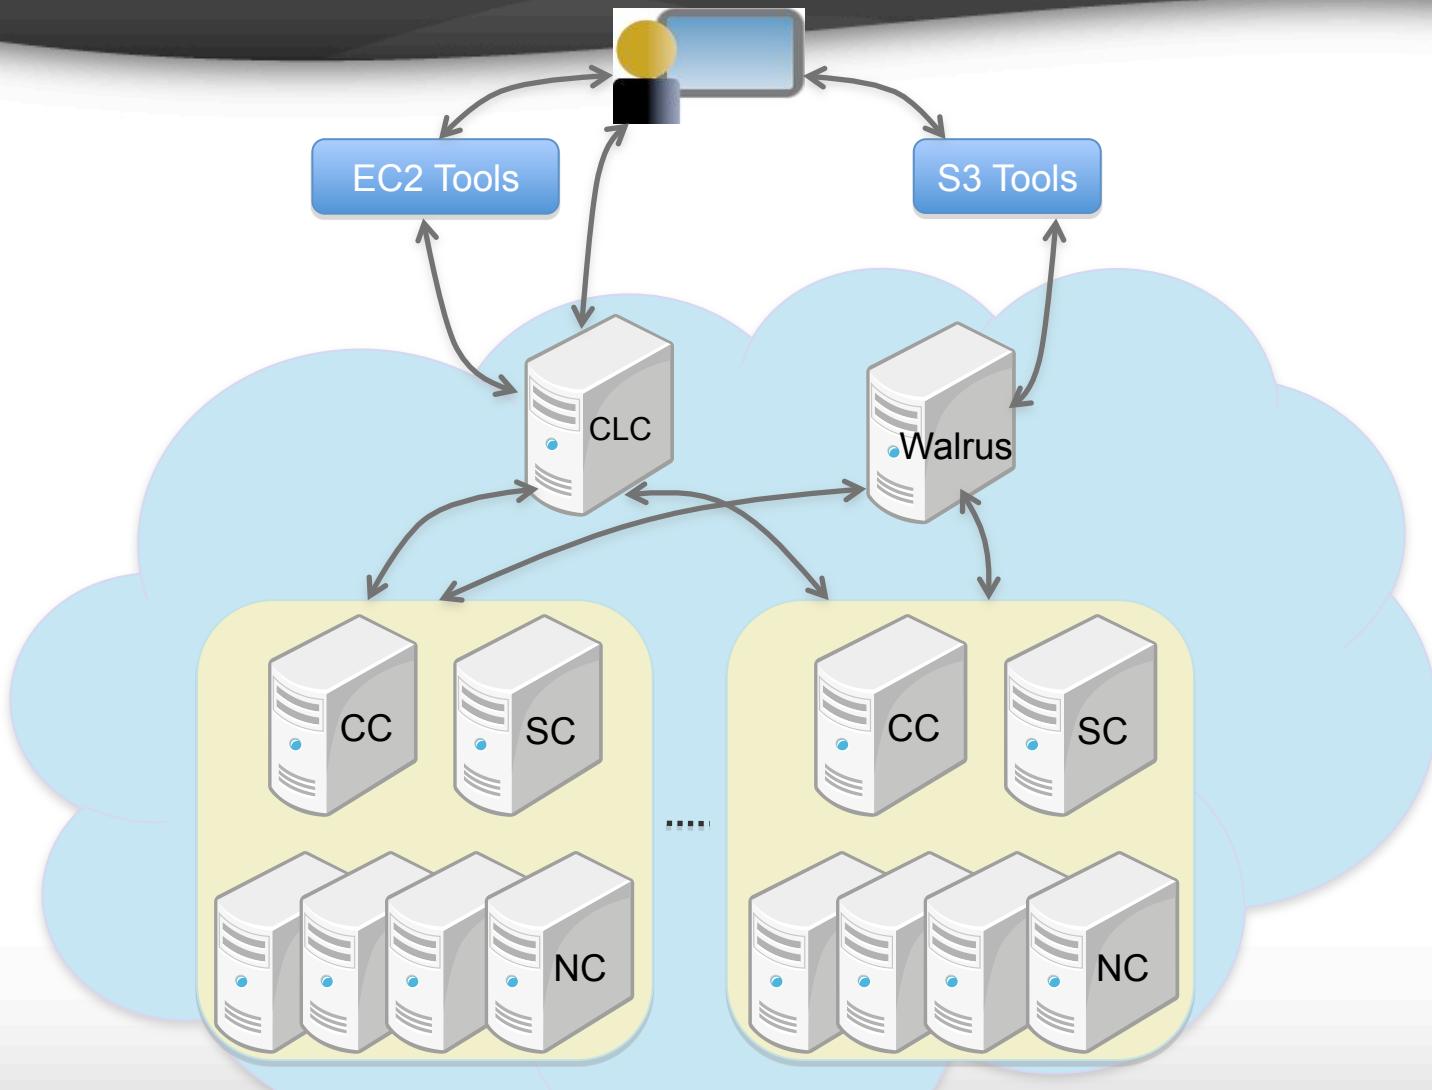
- **Scalability**
 - VM scalability
 - User scalability
 - Storage scalability
 - Node scalability
 - Multi-tenancy
- **User Experience**
 - VM provisioning
 - Request-response times
- **Application Performance**
 - Application I/O characteristics
 - Application architecture
- **Cloud specific**
 - Image management
 - Elastic block storage
 - Security groups
 - Walrus
 - Network configurations

Eucalyptus Architecture



- **Distributed system**
 - Components implemented as web services
- **Components run on top of existing resources**
 - Linux distribution agnostic
 - Ubuntu, RHEL, CentOS, SLES, openSUSE, Debian
 - Hypervisor agnostic
 - Xen, KVM, VMware

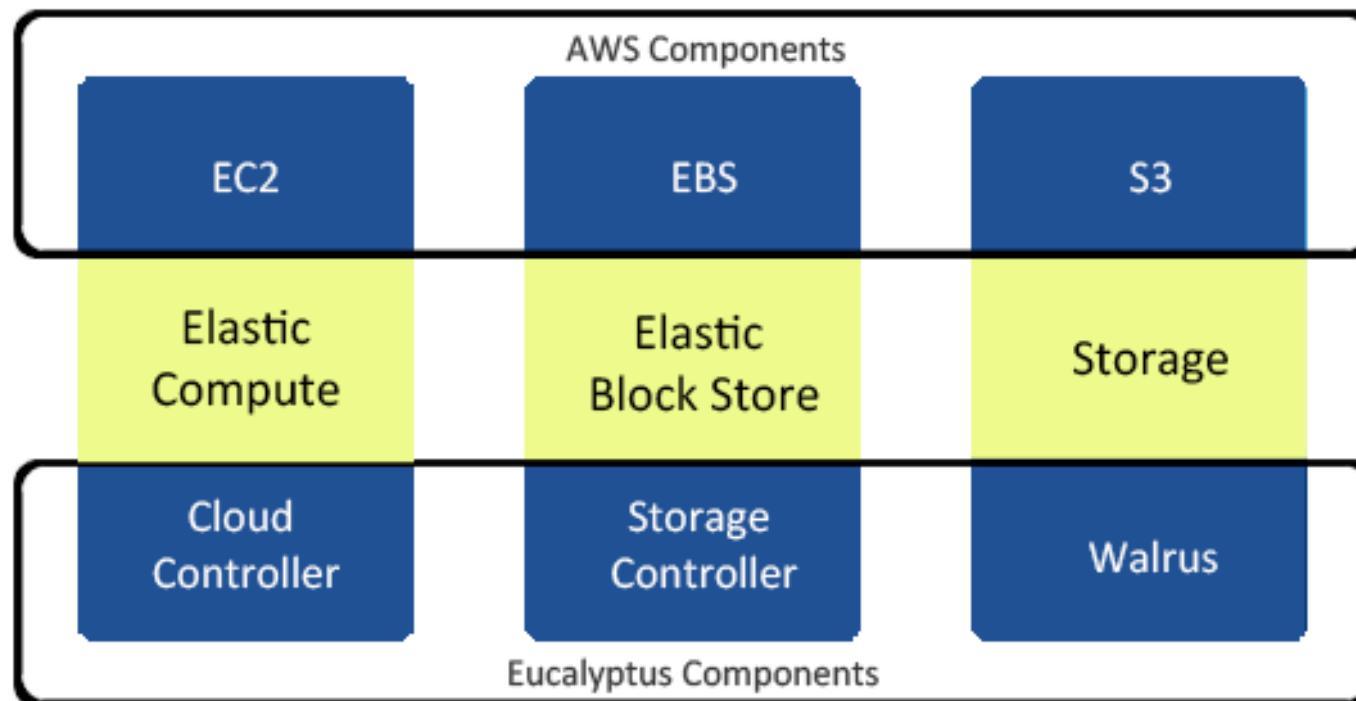
Architecture



Amazon AWS Compatibility



Amazon AWS Public Cloud



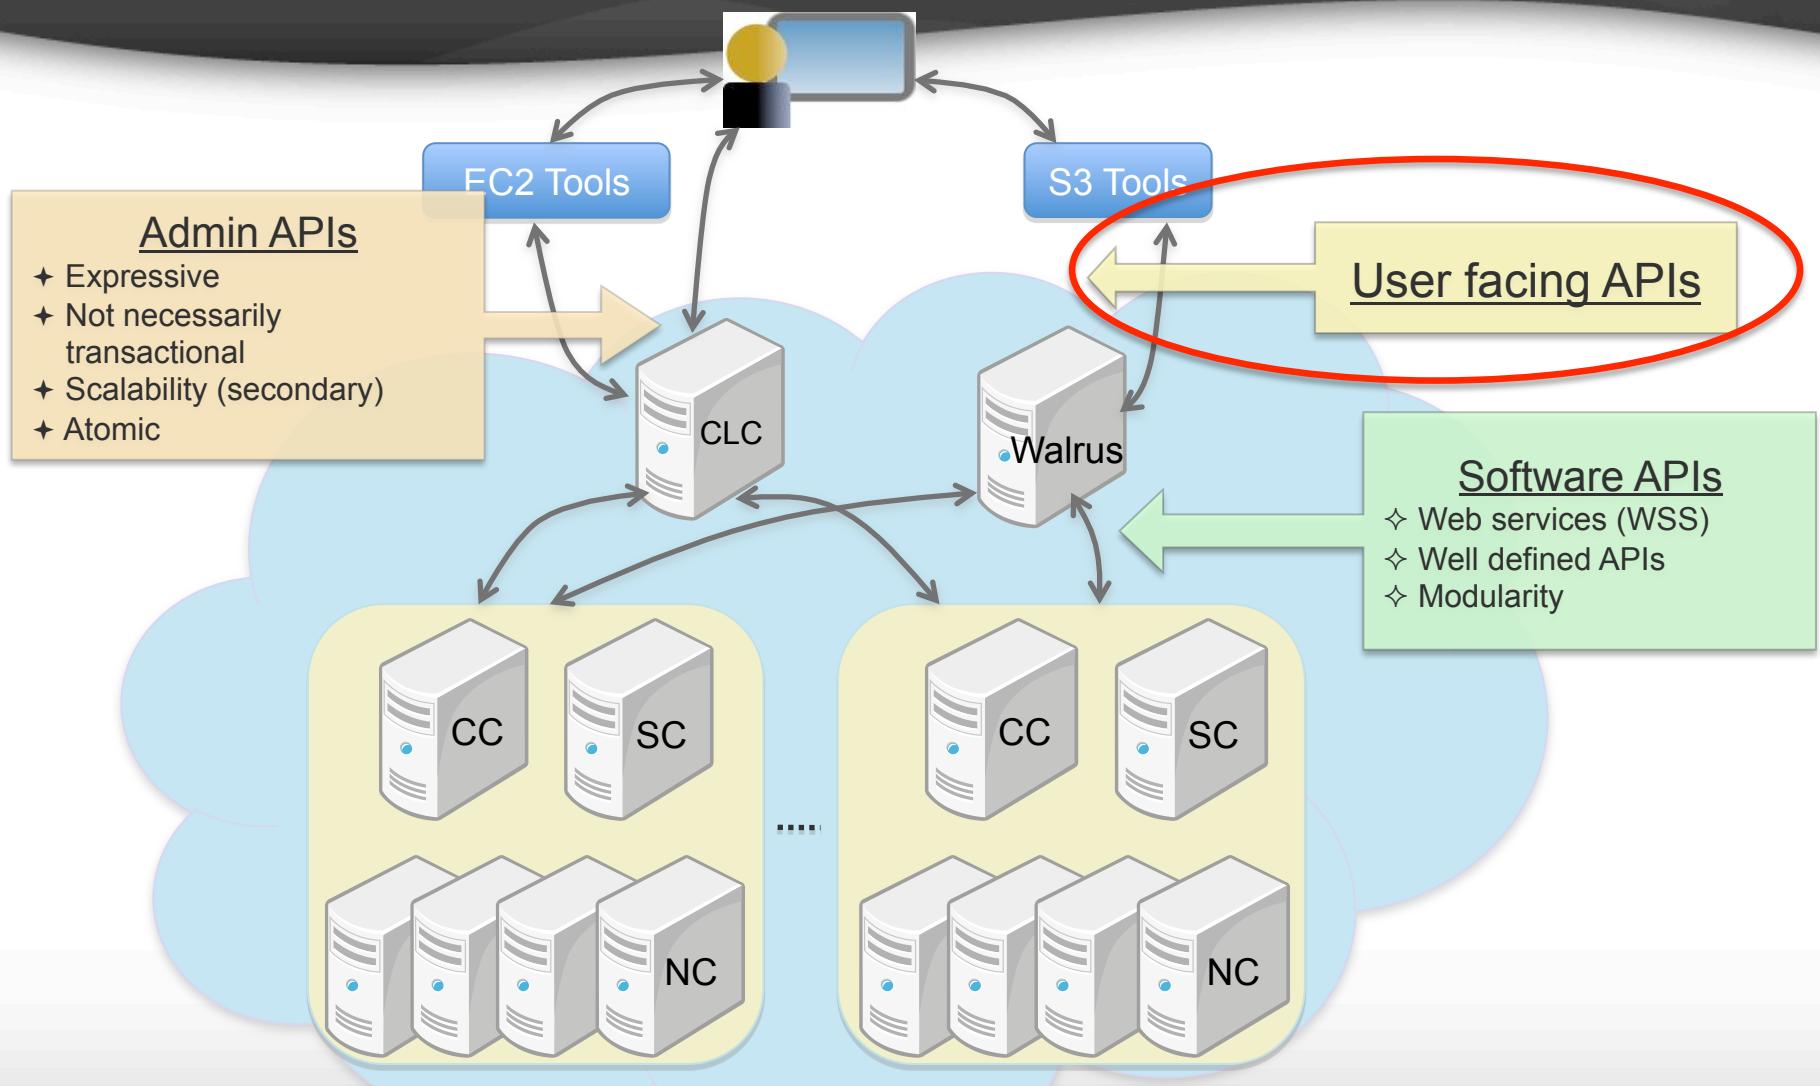
Eucalyptus Private Cloud

Cloud API Requirements



- **Scalability**
 - Resources
 - Users
- **Transactional**
 - Ease of use
 - Failure management
 - Low cost

Distributed Cloud Computing APIs



API Characteristics - I



- **Asynchronous APIs**
 - Deadline driven
 - Client-server agree on a time frame for a response
 - Applications have a defined way to check health/status
 - *User scale*

API Characteristics - II



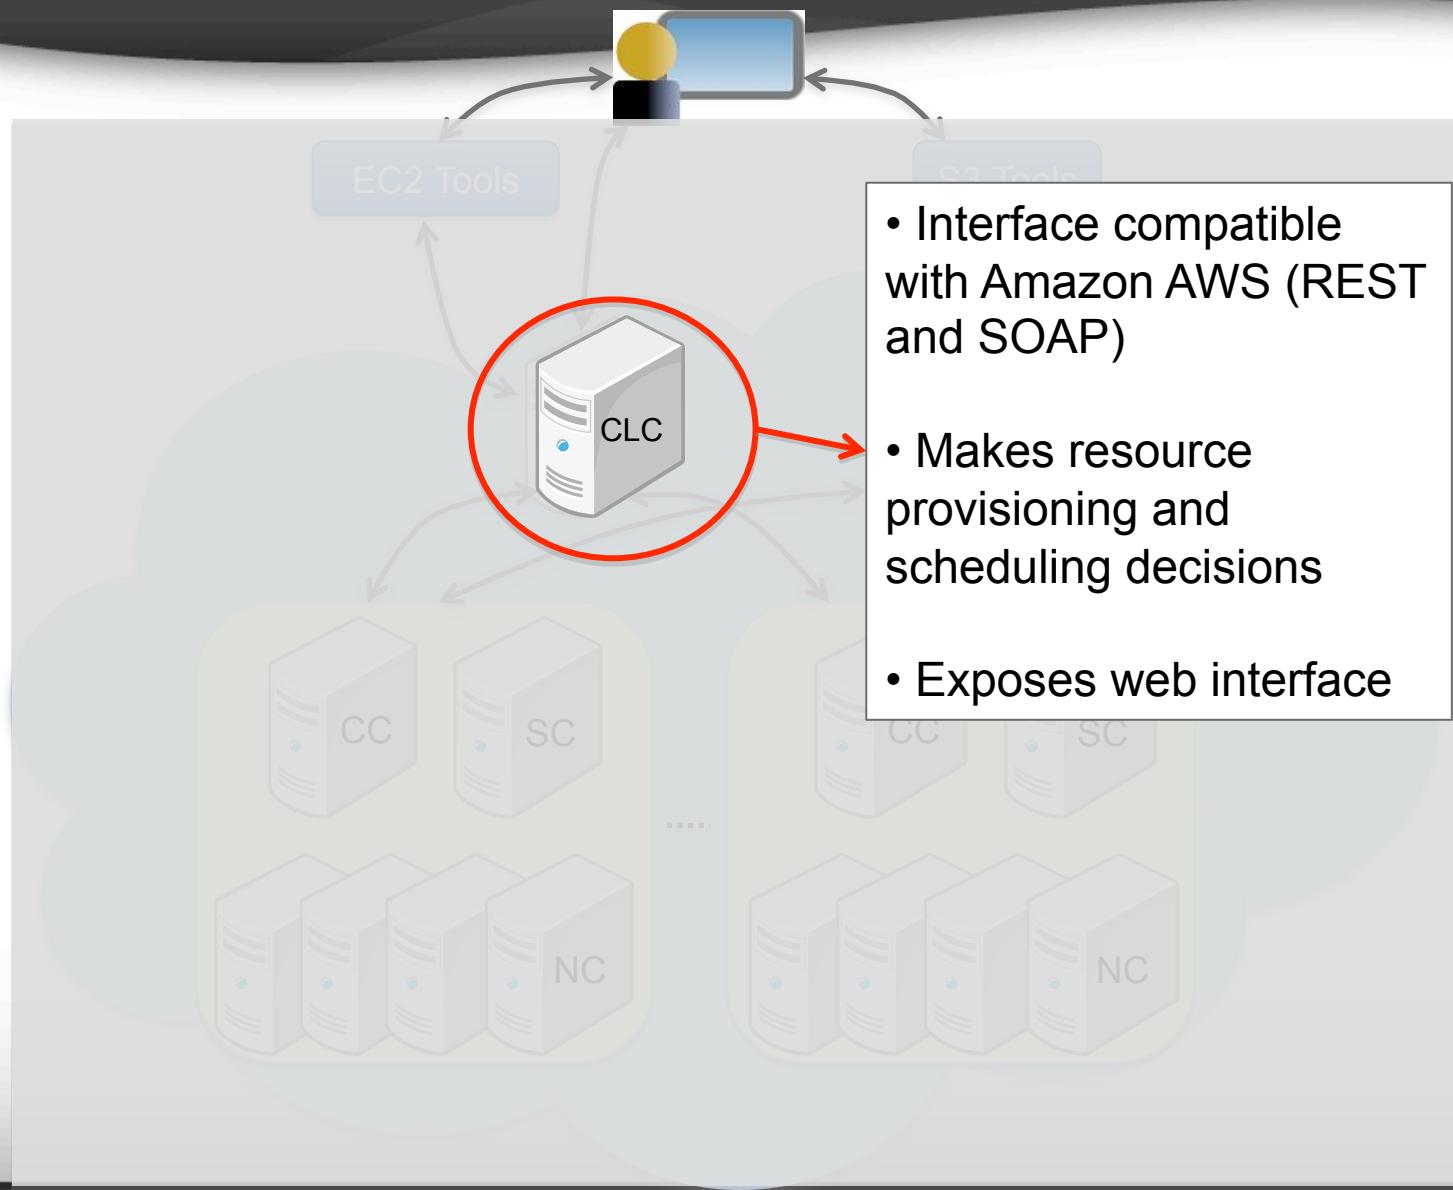
- API avoids aggregate operations
 - Requests refer to “singleton” resources
 - Transactional semantics should scale
 - Simple scheduling
 - No state-locking needed
 - *Resource scale*

API Characteristics - III

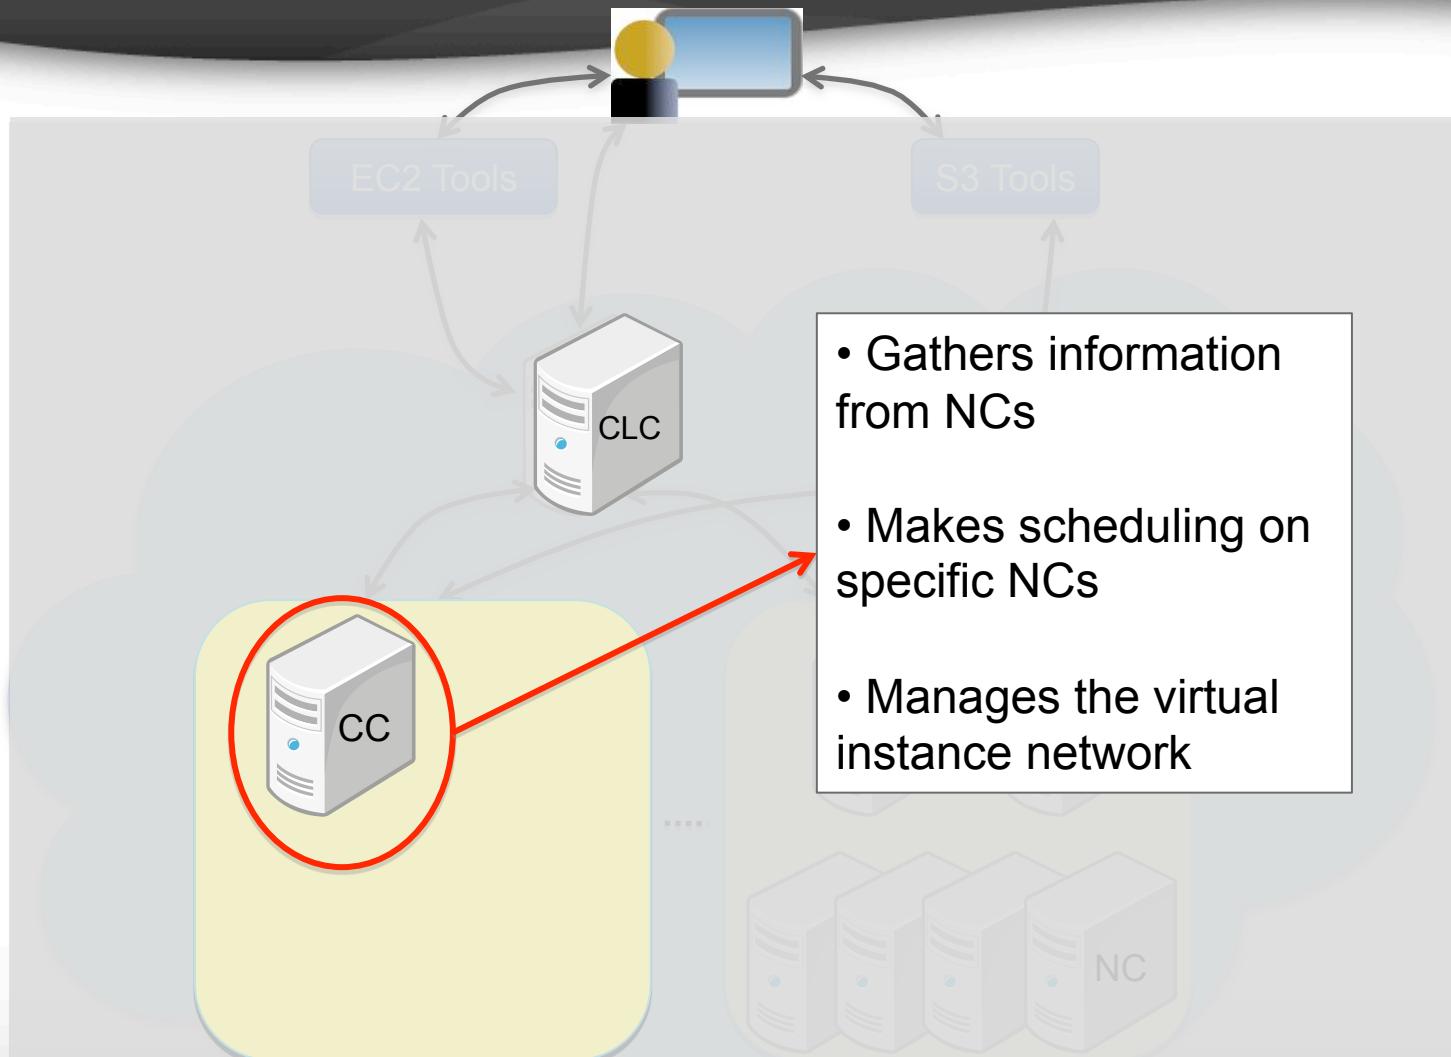


- API mandates scalable persistent abstractions
 - S3
 - Eventual consistency
 - Transactional writes
 - EBS
 - Limited scope to an availability zones
 - No sharing within the cloud semantics
 - Ephemeral
 - Very scalable – 'cos you just destroy them!

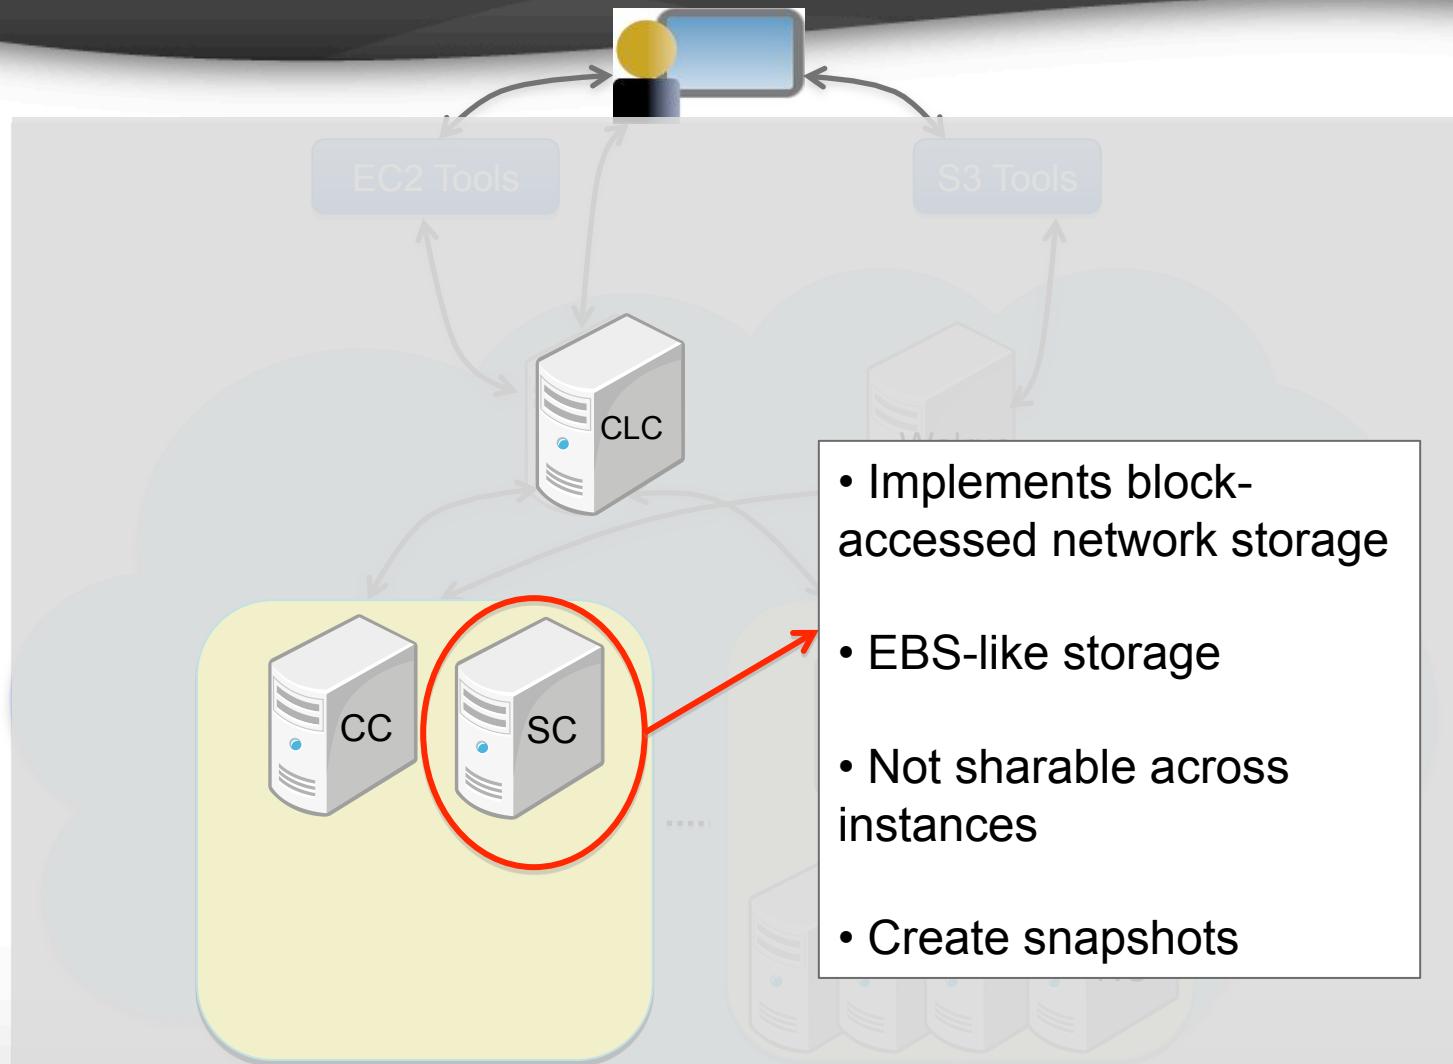
Architecture



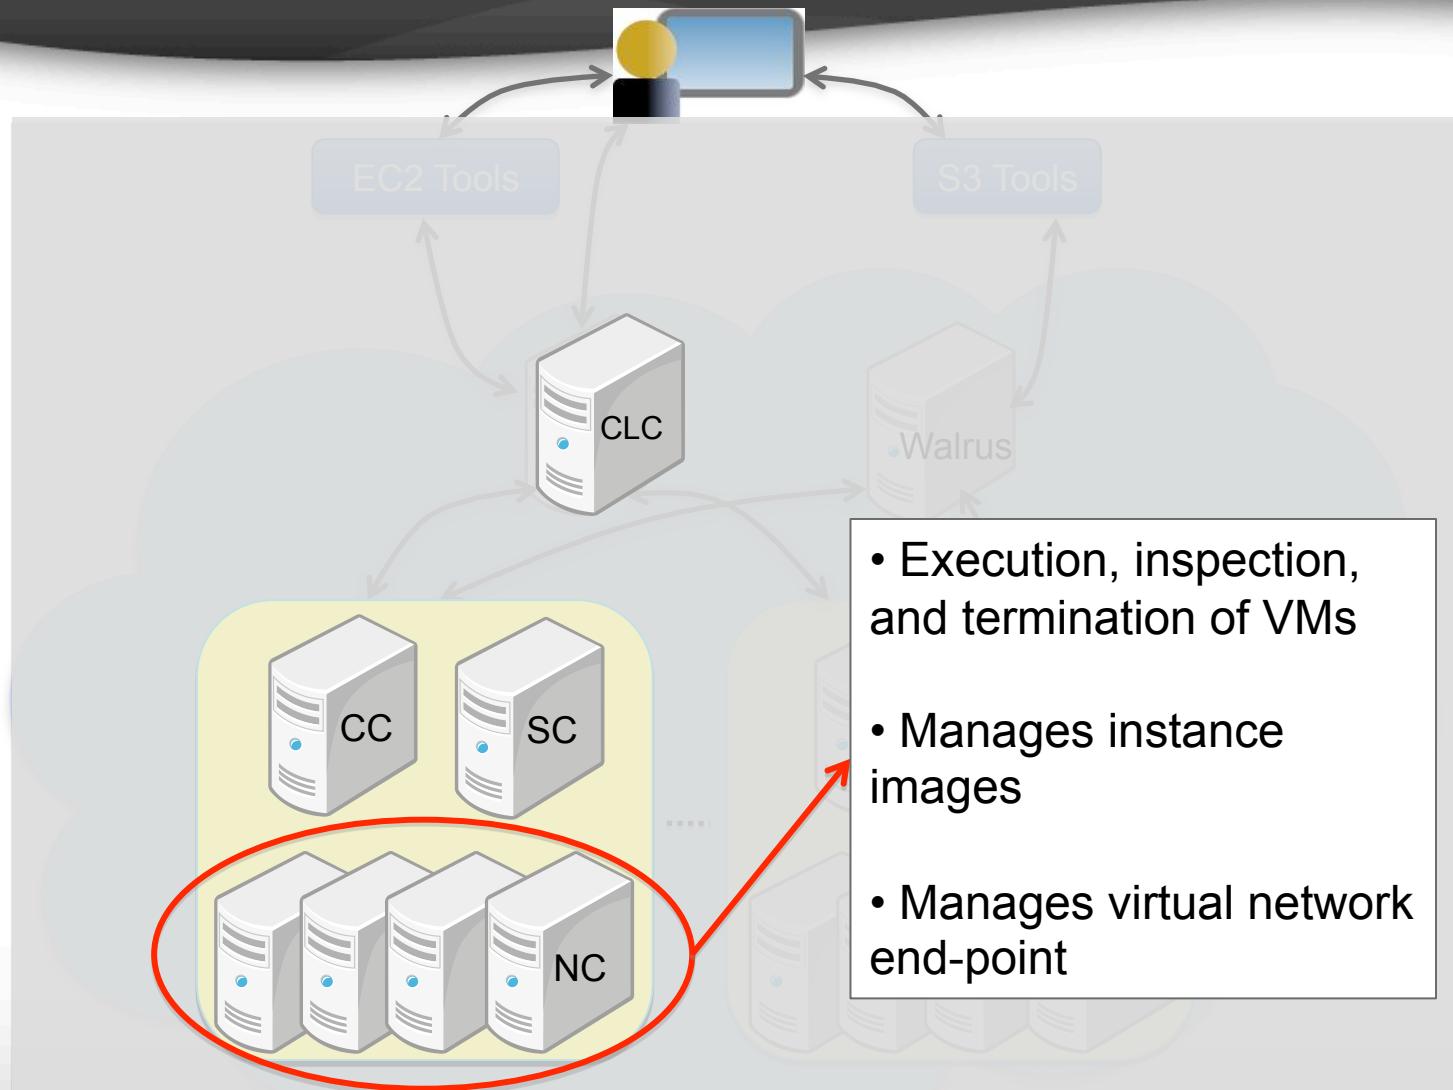
Architecture



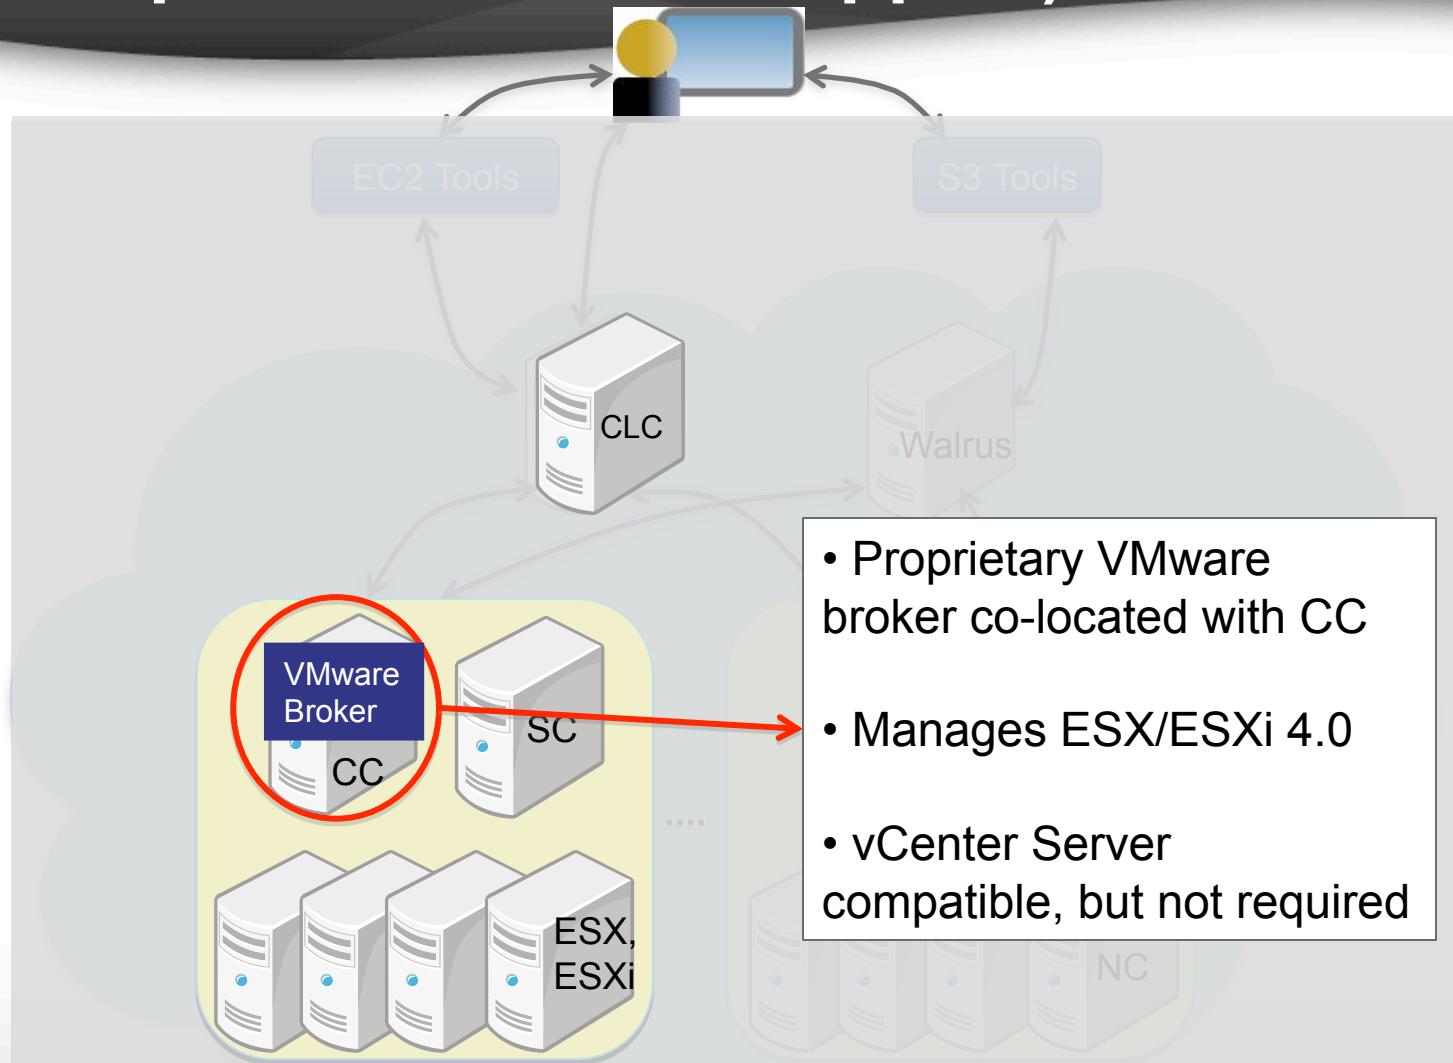
Architecture



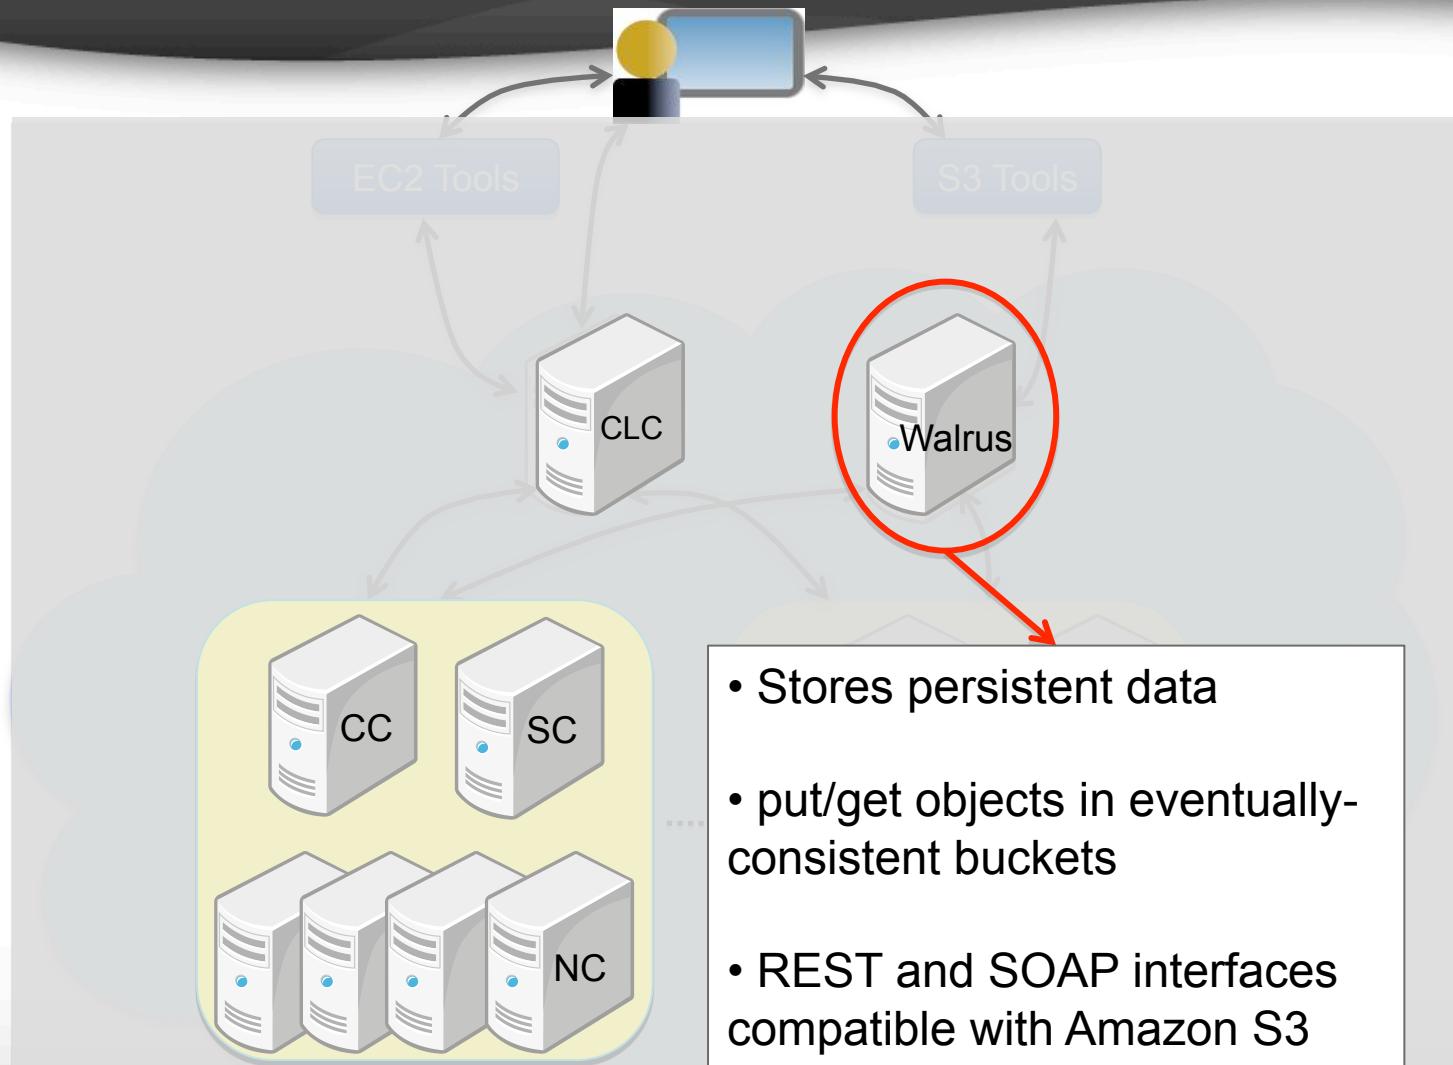
Architecture



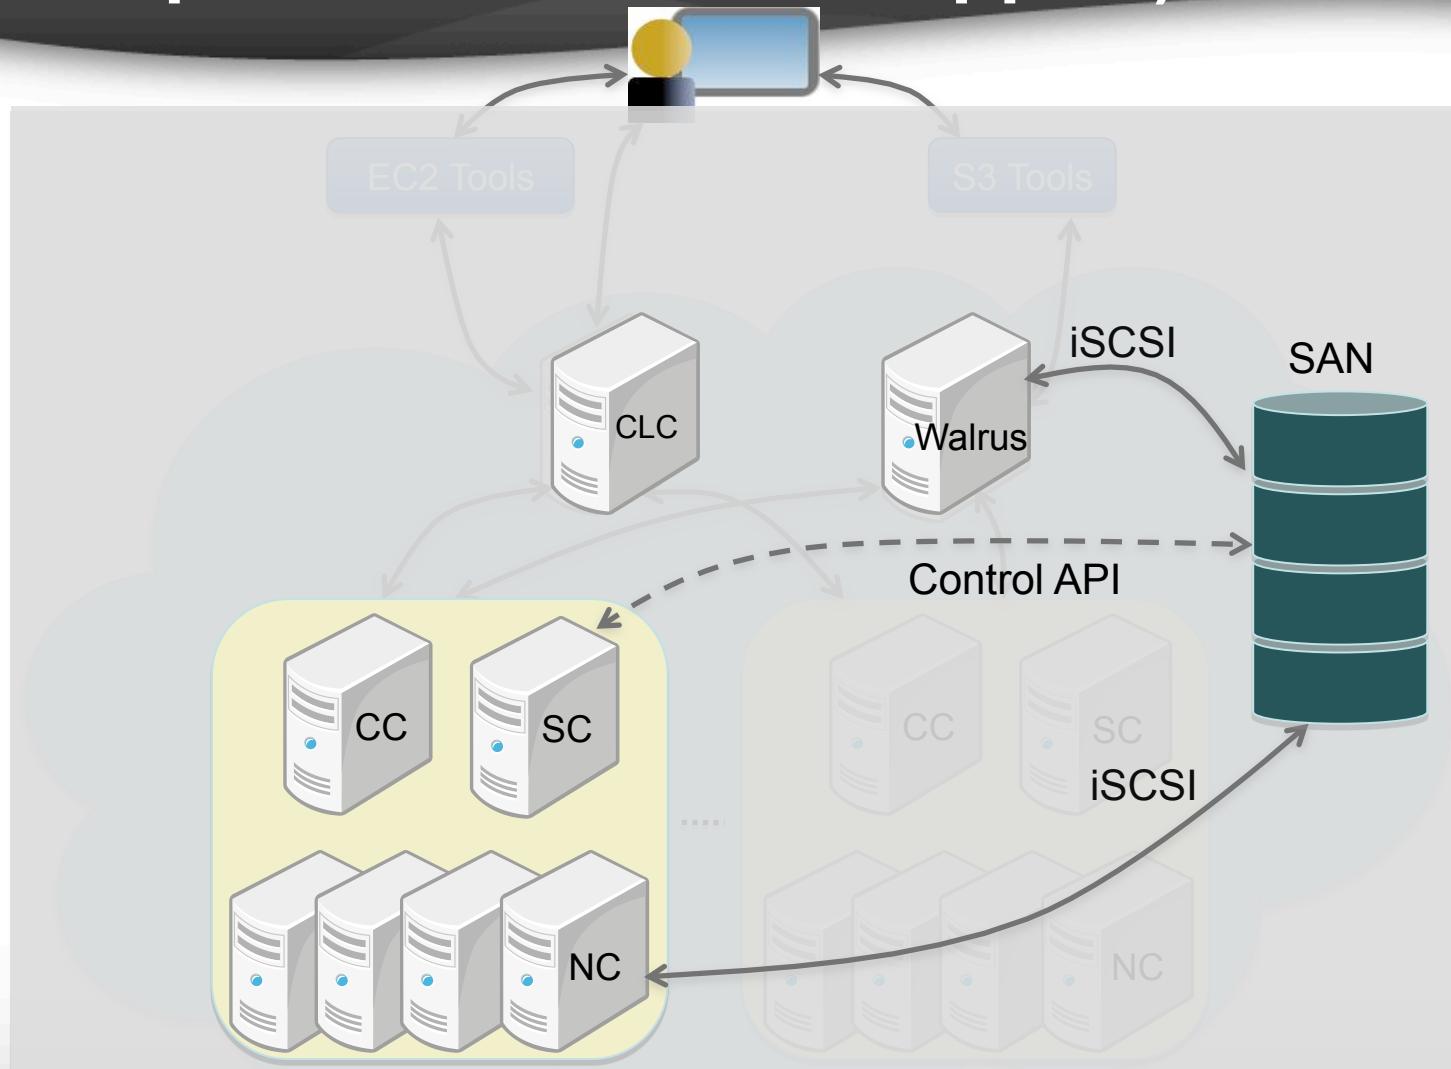
Architecture (Enterprise – VMware support)

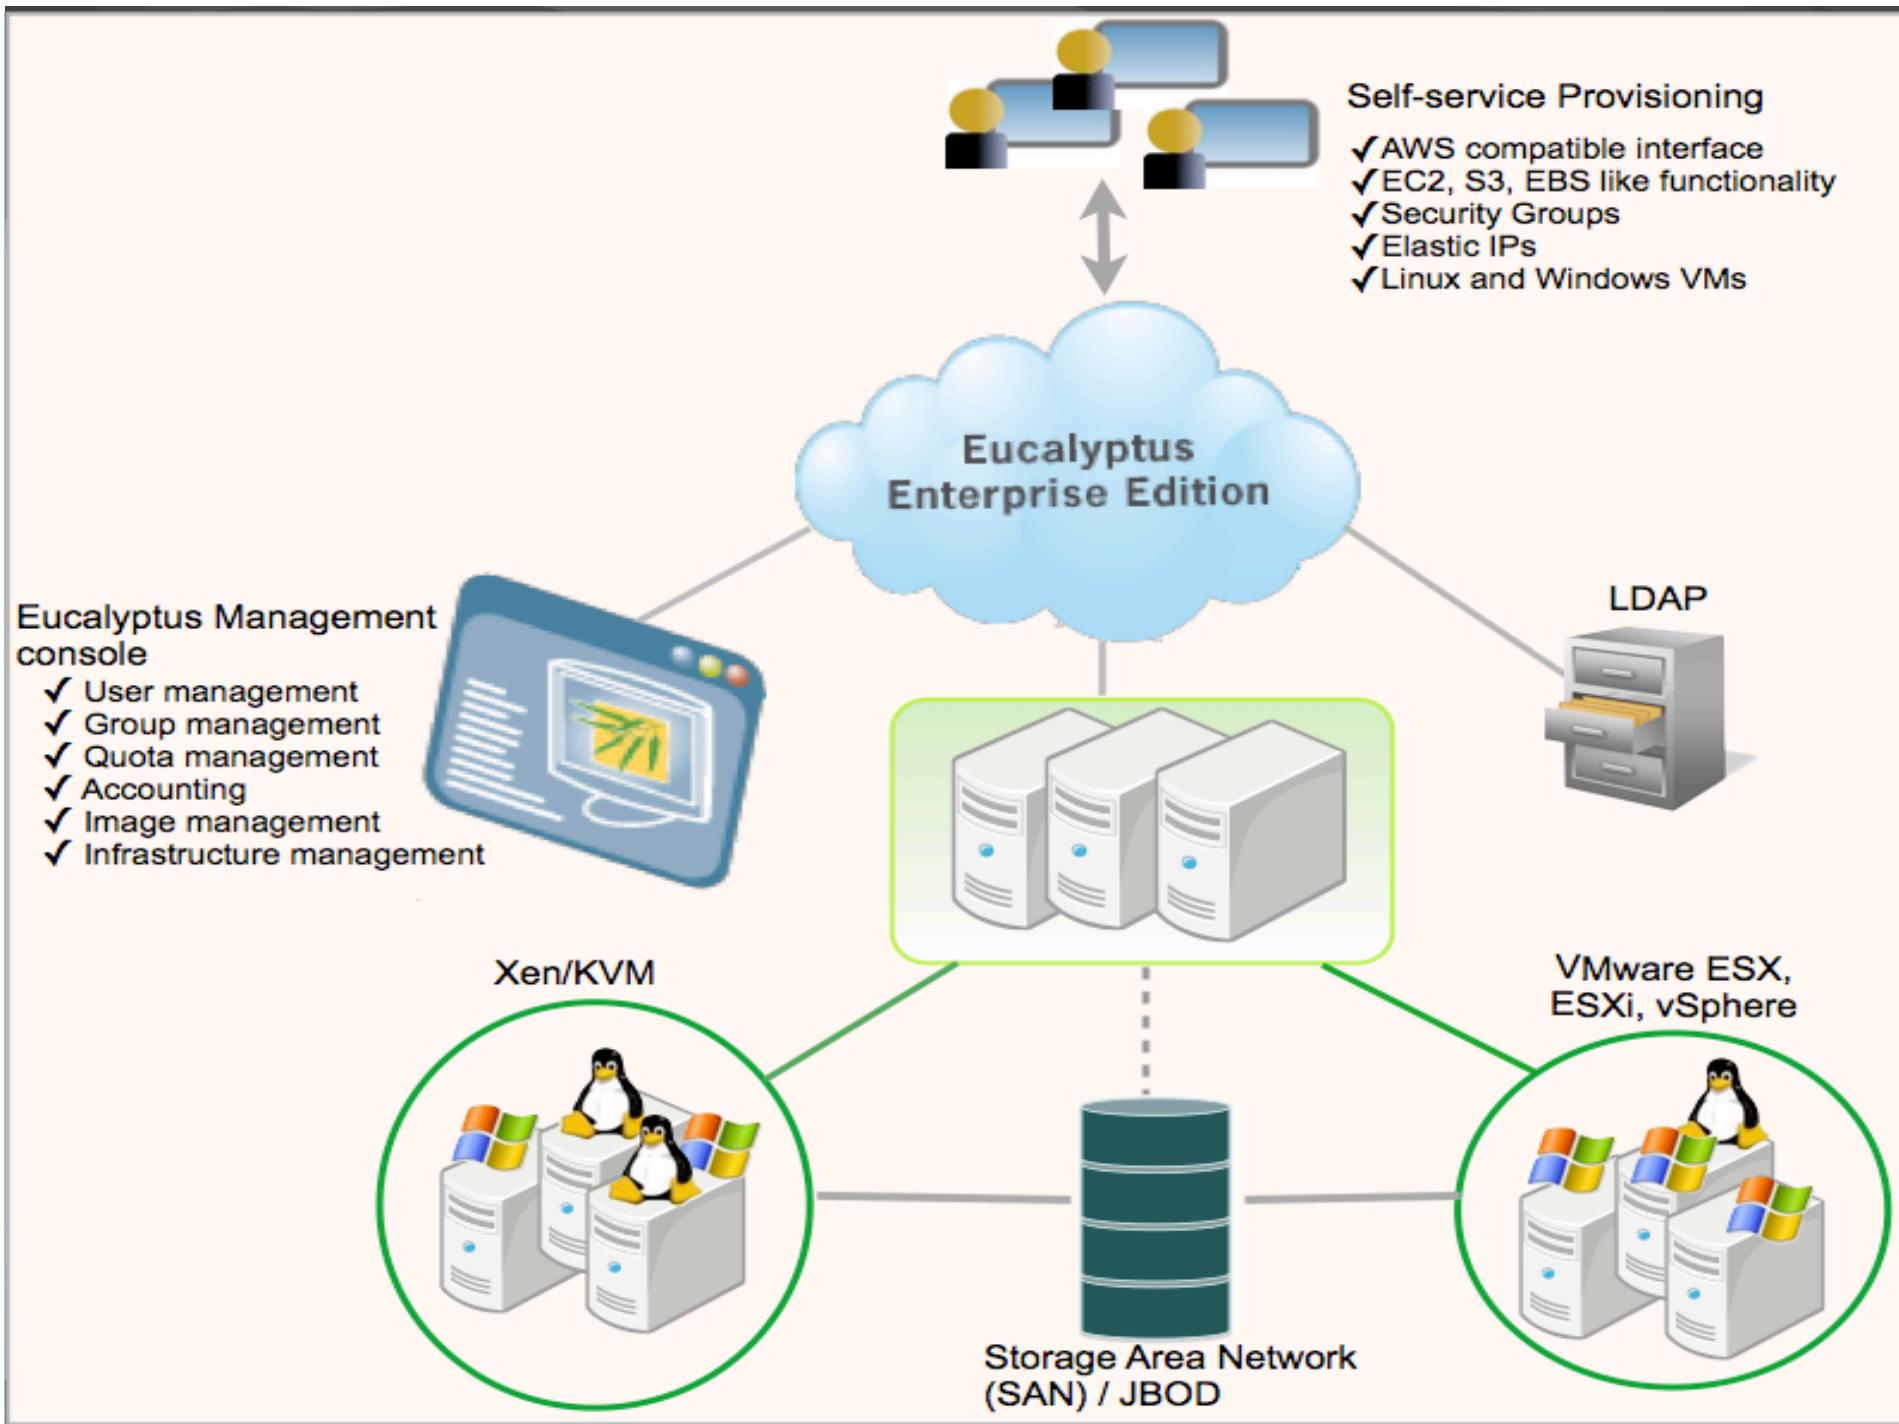


Architecture

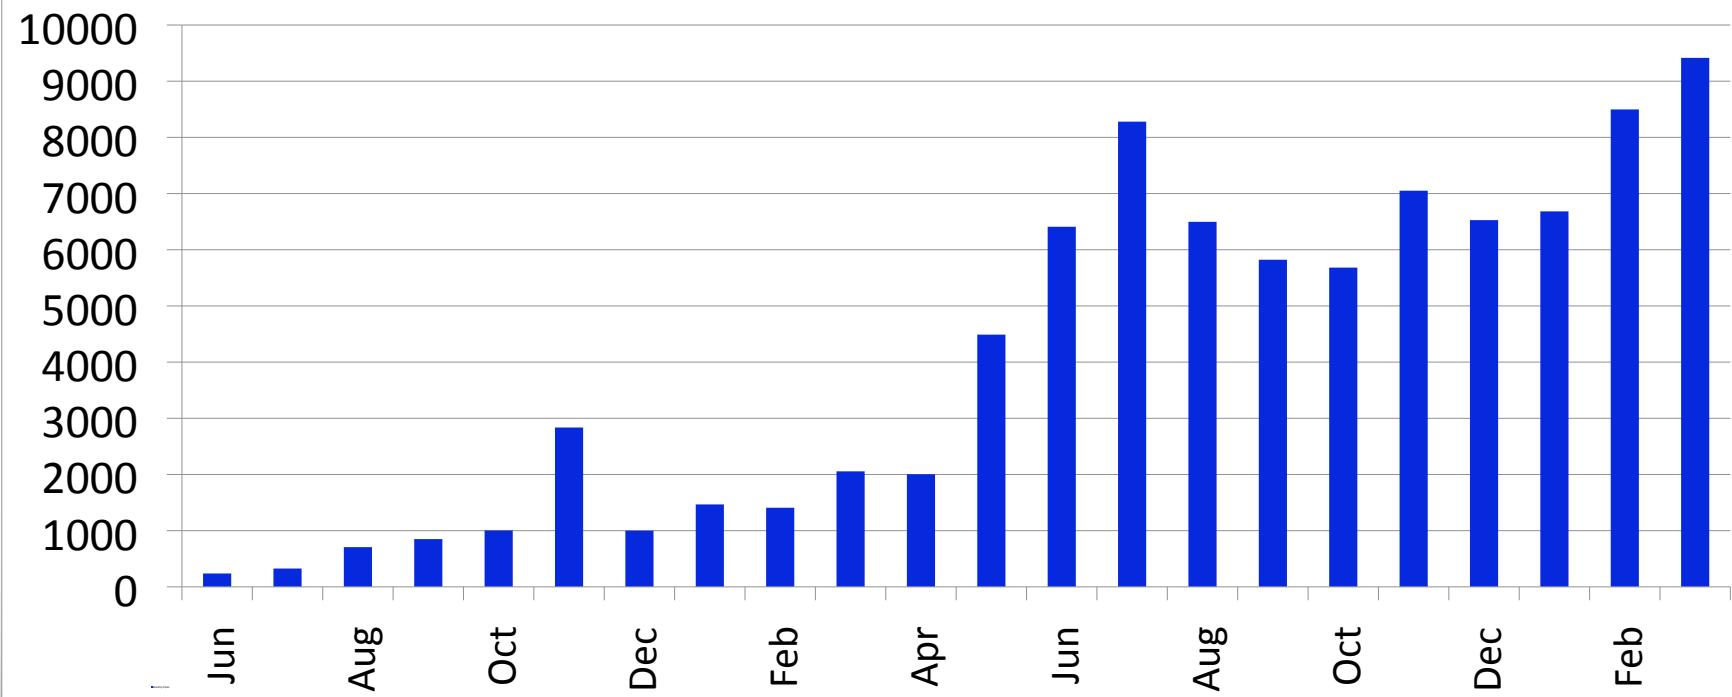


Architecture (Enterprise – iSCSI SAN support)





Downloads (excluding Ubuntu) Monthly Totals



No Eucalyptus in Antarctica (yet)



Open-source Distribution



Linux Distribution: *Ubuntu* and Eucalyptus

- Jaunty Jackalope “Powered by Eucalyptus”
 - April 23, 2009
 - Complete build-from-source
- Karmic Koala
 - October 23, 2009
 - Full-featured Eucalyptus
- Fundamental technology
 - “Ubuntu Enterprise Cloud” ecosystem surrounding Eucalyptus
- 10,000,000 potential downloads
- *Debian* “squeeze”
 - Source release packaging under way
- Packaged for *CentOS*, *OpenSUSE*, *Debian*, and *Ubuntu* as “binary” release as well

A decorative background graphic consisting of several overlapping circles in shades of yellow, orange, and red, creating a cloud-like or network-like pattern.

“Make Eucalyptus the open source reference implementation for cloud computing.”

Simon Wardley (head of cloud strategy), Canonical

Customer Pain Points



- **Self-service procurement of IT resources not available**
- **Automation of resource delivery isn't possible**
- **Procuring machines takes weeks or months**
- **End-users want access to new and legacy resources**
- **More efficient use of resources is required**

Use Cases



- **Testing/Development**
- **Education/Training**
- **Web Services deployment**
- **Quickly create and tear-down IT environments**
- **Hybrid clouds**
 - Cloud bursting
 - Disaster recovery
 - Back up

Thanks!



- Thanks to our original research sponsors...



- ...and to our new commercial friends



www.eucalyptus.com
805-845-8000
shashi@eucalyptus.com